IUM RESEARCH

The Korea Crypto GTM Index

Benchmarks from 25 market entries, 2024–2026. KOL rate cards, the install-to-deposit funnel, listing lead times, event economics, and the timing effects nobody prices in.

Helen, Head of Research · July 2026 · iumlabs.io/blog · Re-cut every two quarters

KEY TAKEAWAYS

- Median Korean KOL pricing runs from 550,000 won per X post at nano tier to 15 million won-plus campaign packages at top tier, and bundle deals close a median 18% below rack rate
- Across attributable CEX-style campaigns, the median funnel converts 1.1% on click-through, 61% install-to-signup, and 27% signup-to-first-deposit, landing blended CAC between $40 and $95 per funded account depending on tier mix
- Timing dominates spend: campaigns launched within two weeks of a major exchange listing saw a median 2.3x engagement multiple on identical budgets
- These are medians from 25 entries, not quotes; treat them as a sanity check against any proposal you receive, ours included

Every founder entering Korea asks the same three questions: what does it cost, what does it convert at, and how long does it take. The honest answer has always been "it depends," because nobody publishes the underlying numbers. Agencies keep rate cards private, exchanges do not discuss lead times, and the case studies that do circulate are cherry-picked wins. This report is our attempt to fix that. We aggregated internal data from the 25 Korea market entries ium Labs executed between H1 2024 and H1 2026, anonymized it, and reduced it to medians. It will not tell you what your campaign will do. It will tell you what the middle of the distribution looks like, which is the number you should negotiate against.

Methodology, Read This First

The dataset covers 25 market entries: 9 exchanges or trading products, 7 infrastructure projects, 5 AI or DePIN networks, and 4 consumer apps. 14 of the 25 had full-funnel attribution from creative to funded account or wallet activation; the rest contribute channel-level data only. All figures are medians unless marked otherwise, because means in this dataset are distorted by two outlier campaigns that outperformed by an order of magnitude. Won figures use a 1,380 KRW/USD convention. Rates reflect what was actually paid after negotiation, not list prices. One bias we cannot remove: this is a client dataset, meaning projects that could afford professional entry, so the floor is likely higher than the true market floor.

25 Market entries in the dataset, executed by ium Labs between H1 2024 and H1 2026, 14 with full-funnel attribution (Source: ium Labs internal)

The KOL Rate Card

Korean crypto KOL pricing is opaque by design, and quoted rates routinely run 30 to 50% above what campaigns actually settle at. Across 214 paid engagements in the dataset, settled medians by tier:

TIER	AUDIENCE	X (TWITTER) POST	YOUTUBE DEDICATED VIDEO	TELEGRAM AMA
Nano	under 10K	550K won (~$400)	rare at this tier	1.2M won
Mid	10-100K	3.2M won (~$2,300)	7.5M won	3.5M won
Top	100K+	12M won (~$8,700)	22M won	8M won

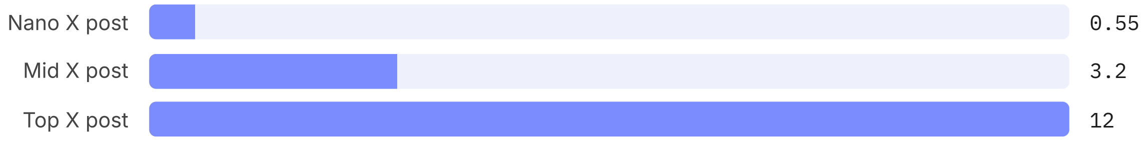

Median settled price per X post by KOL tier, million won. (Source: ium Labs internal, n=214 engagements)

Three patterns worth negotiating around. Bundles beat singles: three-post packages settled a median 18% below the single-post rack rate. YouTube carries a premium of roughly 1.8 to 2x the equivalent X engagement because production is real work in Korea's high-edit-standard scene. And the nano tier is systematically underpriced relative to engagement: nano posts delivered a median engagement rate 2.4x the top tier per won spent, which is why our tier mixes skew nano-heavy for awareness phases.

18% Median discount on three-post KOL bundles versus single-post rack rate (Source: ium Labs internal)

The Funnel Nobody Shows You

For the 14 fully-attributed campaigns, the median funnel from KOL or paid impression to funded account:

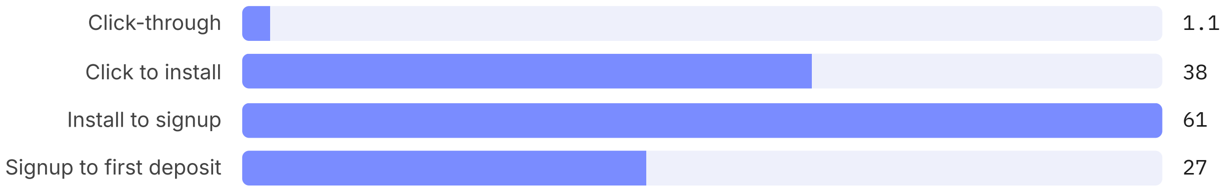

Median conversion at each funnel stage, %. (Source: ium Labs internal, n=14 attributed campaigns)

Multiply it through and 100,000 impressions become roughly 69 funded accounts at the median, which is why impression-count promises are the least useful number in any proposal. Blended CAC per funded account ranged from $40 on nano-heavy mixes to $95 on top-tier-heavy mixes. Paid channels behaved differently: Naver search and display delivered a median CAC 1.7x the KOL blend, but with far better targeting persistence, which is why the stack that won most often was KOL for the spike and Naver SEO plus paid for the tail.

$40-95 Blended CAC per funded account across attributed campaigns, by KOL tier mix (Source: ium Labs internal)

Listing Lead Times

For projects that pursued Korean exchange listings during an engagement, median time from first structured contact to live KRW market: 7.2 months on Upbit and 4.5 months on Bithumb. Both stretched meaningfully after the 2025 compliance tightening; pre-2025 medians were closer to 5 and 3 months. No engagement in the dataset shortened these timelines through spend, and anyone who promises otherwise is describing a different mechanism than the one that exists.

Median months from first structured contact to live KRW market, post-2025 tightening. (Source: ium Labs internal)

Event Economics

Across 38 owned or co-hosted events in the dataset, median attendance was 180 with a range from 60-person closed dinners to 500-plus KBW side events. The useful metric is cost per qualified BD conversation, a contact that progressed to a second meeting: median $85 for owned events, versus roughly $210 equivalent for conference booth presence. Closed-format events with curated lists outperformed open RSVPs on this metric by nearly 3x, which is why headcount is the vanity number and second-meetings-per-won is the real one.

$85 Median cost per qualified BD conversation at owned events, versus ~$210 for conference booths (Source: ium Labs internal, n=38 events)

Community and Search, the Slow Channels

Two slower channels show up in the data with unusually consistent shapes. Telegram communities seeded with coordinated KOL pushes reached 10,000 members in a median 11 weeks versus 26 weeks for organic-only growth, but 90-day retention was statistically identical at around 34%, meaning seeding buys speed, not stickiness. On search, projects that ran the 8-week Naver PR-plus-SEO stack saw median branded-search-volume lift of 210%, and branded Naver search turns out to be the single best leading indicator in the dataset for exchange-side due-diligence attention.

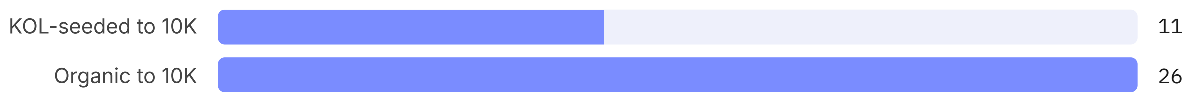

Median weeks for a Telegram community to reach 10,000 members. (Source: ium Labs internal)

The Timing Effect

The largest single variable in the dataset is not budget, tier mix, or creative. It is timing. Campaigns that launched within two weeks of a major catalyst, an exchange listing, a mainnet, a token event, saw a median 2.3x engagement multiple over identical-budget campaigns run in quiet windows. Korean retail attention is event-driven to a degree that global playbooks underestimate, and the practical implication is uncomfortable: a mediocre campaign in the right window beats a great campaign in the wrong one. This is also the strongest argument for sequencing GTM around the listing cycle rather than the marketing calendar.

2.3x Median engagement multiple for campaigns launched within two weeks of a major listing or launch catalyst, versus quiet-window campaigns at equal budget (Source: ium Labs internal)

What Breaks It

Treat every number here with the following discounts. The sample is 25 engagements, large for an agency dataset and tiny for statistics; a single unusual quarter moves these medians. Survivorship bias is real, as failed projects that never engaged professional GTM are absent, and our own failed pitches are absent too. KOL rates drift with the market cycle, and the 2024 bull-leg rates embedded here are likely 15 to 25% above what a 2026 bear-window negotiation would settle at. And attribution in crypto remains partly art: wallet-level attribution on the 14 attributed campaigns is solid for CEX flows and much softer for on-chain actions. We will re-cut this index every two quarters; numbers that stop being true will be replaced, not defended.

Segment Cuts: Who Converts, Who Burns

Medians hide the most useful information in the dataset, which is how differently the four client segments behave. Exchanges and trading products dominate the attributed funnel data and set the CAC benchmarks above, because their conversion event is clean: a funded account either exists or it does not. Infrastructure projects run awareness-weighted mixes where the honest conversion event is developer signups and testnet participation, and their cost per meaningful developer action ran 3 to 5x an exchange's cost per funded account. AI and DePIN projects rode the strongest narrative tailwind in the dataset, roughly 1.4x median engagement on identical spend during narrative windows, but converted to on-chain action worst of the four, because Korean retail trades narratives on exchanges rather than using products. Consumer apps sat in the middle with one structural advantage: KakaoTalk-native referral loops, when they worked, produced the only organic K-factors above 1 we measured.

SEGMENT	N	STRONGEST CHANNEL	WEAKEST CHANNEL	NOTE
Exchanges / trading	9	Nano-KOL bursts + Naver paid	Offline events for retail	Cleanest attribution, CAC benchmarks come from here
Infrastructure	7	Deep-dive content + AMAs	Broad KOL awareness	Developer actions cost 3-5x a funded account
AI / DePIN	5	Narrative-window KOL	Product conversion	Best engagement, worst activation
Consumer apps	4	KakaoTalk referral loops	Paid display	Only segment with organic K-factor above 1

What Changed Against H2 2025

Re-cutting the same metrics against our H2 2025 engagement pool shows three drifts worth pricing in. KOL rates softened roughly 10 to 15% across tiers as the memecoin cooldown cut discretionary campaign demand, with nano rates softening least because supply there is genuinely scarce. Funnel conversion improved at the deposit step, 27% now against roughly 23% in late 2025, which we read as exchange onboarding UX improvements and a higher-intent retail cohort after the app-store delistings of offshore venues. And the timing premium got sharper: the catalyst-window multiple rose from roughly 1.8x to 2.3x, meaning quiet-window campaigns are getting cheaper in absolute terms but even weaker in relative ones. If you are budgeting for H2 2026, the implication is uncomfortable but simple: the calendar is now a bigger lever than the creative.

The Quarter Ahead

Three scheduled forces shape Q4 2026 GTM economics, all previously covered in this library. The 2027 gains tax arrives January 1, and the behavioral fallout we mapped, pre-deadline realization and offshore drift, should pull retail attention forward into Q4, making it the densest catalyst window of the cycle. The security-token framework takes effect January 2027, which will pull institutional narratives into year-end. And the corporate-access phases keep grinding forward, adding a B2B lane to what has been a purely retail market. Our operating read: Q4 2026 is the strongest launch window Korea has offered since early 2024, and the KOL rate softening documented above will not survive it.

The Search Layer, Underpriced and Slow

Every benchmark above decays; search compounds, which is why we keep it in the index despite it being the least glamorous line. Across engagements that ran the 8-week Naver stack, PR placement into Naver-indexed outlets, blog-network seeding, and branded keyword buys, the median branded-search lift was 210%, and that lift persisted at roughly 60% of peak six months later with zero incremental spend. Compare that to KOL bursts, where engagement returns to baseline within two weeks of the last post, and the budget logic writes itself: KOL buys the spike that makes people search, and the Naver layer decides what they find when they do. Two practical medians for planning: a defensible branded-keyword position on Naver cost 4 to 9 million won per month depending on category competition, and the PR component only moved search when at least two of the placements landed in outlets Naver surfaces in its news vertical, which is a property of the outlet list, not the story. The failure mode we saw twice and now screen for: teams that bought the spike, skipped the layer, and watched competitors' content rank for their own brand name during their highest-attention week.

210% Median branded Naver search-volume lift after the 8-week PR-plus-SEO stack, persisting at ~60% of peak six months later (Source: ium Labs internal)

Reading a Proposal With This Index, a Worked Example

Here is how the index earns its keep in a real negotiation. Take a composite proposal we saw this spring: 120 million won for a six-week awareness campaign, promising 10 million impressions via 15 KOLs, 20,000 installs, and an exchange listing "accelerated through relationships." Line one: 15 KOLs at a 120M-won blend implies 8M won per engagement, which sits at our top-tier median, so either every name on the list is genuinely top-tier, ask for the list, or the blend is padded. Line two: 10 million impressions converting to 20,000 installs implies a 0.2% impression-to-install rate, which our funnel medians say is optimistic by roughly 2x for a KRW-focused mix, so ask which markets the traffic comes from. Line three: 20,000 installs at the median install-to-signup and signup-to-deposit rates yields about 3,300 funded accounts, putting proposed CAC near 36,000 won, plausible, but only if the funnel holds, so tie at least 30% of fees to the deposit event, not the impression count. Line four: no spend accelerates a listing timeline, our medians are 7.2 and 4.5 months and nothing in 25 engagements shortened them, so strike the clause and the fee attached to it. Ten minutes with the index turned a 120M-won ask into three pointed questions and one deleted line item, which is precisely the leverage we intend it to give you, including against us.

Appendix: How We Collected It

Attribution ran through UTM-tagged links, exchange referral codes, and where clients permitted it, postback data from their onboarding funnels; on-chain attribution used tagged deposit addresses and is materially weaker, which is why only 14 of 25 engagements clear our full-funnel bar. Rates are settled invoice amounts, not quotes. Engagement medians are per-campaign, not per-post, to stop single viral posts from distorting tiers. The H2 2025 comparison pool uses identical definitions. We publish medians and ranges rather than means throughout, and we will re-cut this index every two quarters with the same methodology so the numbers stay comparable. Corrections, challenges, and requests for cuts we did not publish: research@iumlabs.io.

How to Use This

If you are evaluating a Korea entry proposal, ours or anyone's, benchmark it here. A KOL quote dramatically above the tier medians needs a reason. A funnel projection above the stage medians needs a mechanism. A listing timeline below the medians needs a miracle. The gap between what you are being sold and what the middle of the distribution actually does is the most useful number in this report.

Sources

ium Labs internal campaign dataset, 25 Korea market entries, H1 2024 to H1 2026, aggregated and anonymized. Exchange lead-time figures reflect engagements where ium Labs advised on listing readiness; they are observational, not exchange-confirmed. FX convention 1,380 KRW/USD. Public context figures referenced from DAXA disclosures and Korea Financial Intelligence Unit reporting. This report contains no client-identifiable information. Benchmarks are directional and do not constitute a quote or guarantee.

